Jainish Parikh

jainishparikh.github.io | 408-828-4888 | jainishjparikh@gmail.com | linkedin.com/in/jainish-parikh | github.com/jainishparikh

Backend engineer focused on building reliable, high-performance data systems. Primary owner of 6sense's query-generation platform used across core products, driving major improvements in throughput, latency, and efficiency. Experienced in distributed data processing, **performance optimization**, and **end-to-end service ownership**

Experience

6sense Insights Inc, San Francisco, CA, US   *Dec 2021 - Present*
Software Engineer

- **Owner** of 6sense's core data-access layer and dynamic query-generation framework, responsible for performance, reliability, and long-term evolution.
- Designed and evolved a query-generation engine converting declarative JSON specifications into optimized executable queries, powering **4 core products**.
- Owned operations of the company's most compute-intensive workflow — a nightly job executing **200K+ complex customer queries** — ensuring correctness and availability.
- Delivered major performance improvements: **1.5× faster average runtimes**, **6× improvements for ~20% of heavy (>20s) queries**, and **92% lower memory usage**, significantly increasing system efficiency and reducing costs.
- Acted as the domain expert for query-generation, mentoring engineers, writing internal guidelines, and helping multiple teams adopt the framework effectively

Juniper Networks, Sunnyvale, CA, US   *July 2021 - Dec 2021*
Software Engineer

- Created **dashboards, alerts, and metrics pipelines** in SignalFx and AWS CloudWatch via Terraform, improving observability.

Intel Corporation, Santa Clara, CA, US   *September 2020 – December 2020*
Software Engineer Intern

- Built **CI pipelines in GitLab** to automate Docker container builds, reducing job-queue length by >50%.
- Automated dynamic package installation and image storage in Artifactory using Shell scripts.

Projects

NoSQL peer-to-peer Database (Java, REST, Docker, CAP, Distributed systems)

- Designed and built a peer-to-peer distributed NoSQL database with **high availability**, **partition tolerance**, and **eventual consistency**, inspired by CAP theorem tradeoffs.
- Implemented multi-node replication where nodes **gossip state every 2 seconds**, maintaining cluster-wide convergence.
- Used **vector clocks** to handle CREATE/UPDATE/DELETE conflicts across replicas, ensuring correct causal ordering.
- Exposed a set of REST APIs for data access, replica synchronization, and consistency management across nodes.

Yelp Prototype (MongoDB, ReactJS, NodeJS, REST, GraphQL ,Docker, AWS, Kafka, Redis)

- Engineered a **modular full-stack application** featuring event-driven backend services for search, orders, reviews, and geolocation functionality.
- Built scalable API services backed by **Kafka-driven asynchronous workflows**, allowing the system to handle high concurrency without degradation.
- Added **Redis caching layers** to accelerate query performance for search and restaurant details.

Education

M.S. in Software Engineering	San Jose State University, USA	*August 2019 - May 2021*
B.E. in Information Technology	L.D. College of Engineering, India	*August 2015 -May 2019*

Skills

- **Languages:** Python, Java
- **Backend & Data:** REST, GraphQL, Kafka, Redis, SingleStore, MongoDB, Query Engines, Distributed Systems Concepts, Performance Optimization
- **Cloud & Infra:** AWS, Docker, Kubernetes, Terraform, CI/CD
- **Frontend (supporting):** React, TypeScript, JavaScript